

# RL Algorithms comparison | Yessmine

## Requirements:

offline

easy to implement

focus on discrete (Reward Function)

## Model-Free RL:

- Involves methods that learn to act without needing a model that predicts the environment's behavior.
- Contrasts with model-based RL, which builds a model of the environment to simulate and plan actions.

## On-Policy RL:

- **Learning Approach:** Improves the policy that is currently being used.
- **Examples:** SARSA, on-policy actor-critic methods.
- **Advantages:** Simplicity, direct policy improvement.
- **Disadvantages:** Less sample-efficient, challenging exploration-exploitation trade-off.

## Off-Policy RL:

- **Learning Approach:** Improves a different policy from the one used to generate data.
- **Examples:** Q-learning, DQN.
- **Advantages:** More sample-efficient, flexible exploration.
- **Disadvantages:** More complex, potential stability issues.

## Zied's Research Outcome

Algorithms that are useful for offline learning with Discrete action and observation spaces, that are Model free , and off policy (exploration),

## offline DQN, QR-DQN and Batch-Q Learning

\*

QR-DQN usually shows better results but is a bit more complex to implement.

Implementation :

-Like online algorithms but has to be adapted (still to research but seems doable)

## Yin's Research Outcome

### BCQ

#### ▼ ~~Tool D4RL: Algo evaluation Tool~~

Rich Baselines: D4RL provides a comprehensive set of baselines, including common offline algorithms like BCQ, BEAR, BRAC, and others.

## Evaluation

- BCQ is an advanced offline RL algorithm specifically designed to mitigate issues inherent in BQ, such as overestimation and extrapolation errors.

BQ as subcategory of DQN.

- If you have a relatively simple problem and want a straightforward implementation, you might start with Offline DQN or Batch Q-Learning. → Also the easiest
  - Offline DQN: can provide a solid foundation in RL concepts, including value-based methods, experience replay, and neural network approximation.
  - **Batch Q-Learning:** can help you understand the basics of Q-learning algorithms and how they operate without the added complexity of deep neural networks.
- If capturing uncertainty in Q-values is crucial for your application, QR-DQN(**Quantile Regression DQN**) could be a good choice.

- If you're dealing with complex environments and want a method explicitly designed for offline RL with improved exploration strategies, BCQ (**Batch-Constrained Q-learning**) might be worth exploring.

→ BQ, BCQ or QR-DQN.

## Decision

BQ

## Q-Learning: (steps)

1. **Initialize** the Q-function arbitrarily (e.g., all zeros).
2. **Observe** the current state  $s$ .
3. **Select** an action  $a$  using an exploration policy (e.g.,  $\epsilon$ -greedy).
4. **Take** the action and observe the next state  $s'$  and the reward  $R_{t+1}$ .
5. **Update** the Q-value using the Bellman equation:

$$Q(s,a) \leftarrow Q(s,a) + \alpha [R_{t+1} + \gamma \max_{a'} Q(s',a') - Q(s,a)]$$

where  $\alpha$  is the learning rate.

6. **Repeat** steps 2-5 for each time step and episode until convergence.